

Introduction to ANNs and Limitations in Image Processing

Course:

INFO-6152 Deep Learning with Tensorflow & Keras 2



Developed by:
Mohammad Noorchenarboo

September 9, 2025

Current Section

1 Course Introduction and Marking Scheme

2 Artificial Neural Networks (ANNs)

3 Using ANNs for Image Data (28×28)

4 Loss of Spatial Structure in ANNs

5 High Number of Parameters in ANNs

6 Lack of Translation Invariance in ANNs

Course Goal

This course **INFO-6152 Deep Learning with TensorFlow & Keras 2** focuses on advanced concepts, architectures, and applications of deep learning. We build upon the foundations of TensorFlow and Keras to explore modern models and their deployment in real-world contexts.

Main Objectives

- Extend understanding of deep learning beyond the basics
- Explore advanced architectures: CNNs, RNNs, LSTMs, GRUs, Attention, Transformers, GANs
- Learn strategies for improving model performance and generalization
- Apply transfer learning and fine-tuning in TensorFlow/Keras
- Understand generative models and their applications
- Gain practical skills in deploying and scaling deep learning systems

Course Goal

Outcome

By the end of this course, students will be able to:

- **Design and train** advanced neural network architectures
- **Implement and fine-tune** models for vision, sequence, and generative tasks
- **Apply** TensorFlow/Keras in solving real-world problems
- **Deploy and optimize** models for efficiency in production settings
- **Critically evaluate** deep learning solutions and trade-offs

Course Assessment Overview

Marking Scheme (100 points total):

Component	Weight (%)
Quizzes	20
In-class Activity	10
Assignments	10
Final Exam	30
Final Project	30

In-class Activities (10%)

Rules

- Students respond to questions after lectures. Each clear and complete answer earns **5 points**.
- Across the course, each student must answer exactly two times.
- One answer must occur **before the Reading Week**, and one **after**.
- Students may volunteer; otherwise, the instructor will select randomly.

Quizzes (20%)

Structure

- Multiple-choice format
- No cheat sheets allowed
- Each quiz has **10 questions**, each worth **0.5 points**
- Total = **5 points per quiz**

Assignments (10%)

Description

- Short, coding-based or written tasks on deep learning topics
- Students complete **two practical in-class activities** applying TensorFlow/Keras
- Each assignment must be **completed and submitted during class hours**
- Activities reinforce advanced deep learning concepts such as transfer learning, regularization, and generative models

Final Exam (30%)

Format

- Multiple-choice questions
- **30 questions**, each worth **1 point**
- One page (two-sided) cheat sheet allowed

Final Project (30%)

Requirements

- Students work **independently** (no groups)
- Must apply TensorFlow and **advanced deep learning techniques** on a real dataset
- Includes both implementation and a short presentation
- Emphasis on creativity, correctness, application, and clarity

Important Notes

Reminder

- All dates for quizzes, assignments, and exams are specified in the course outline.

Current Section

1 Course Introduction and Marking Scheme

2 Artificial Neural Networks (ANNs)

3 Using ANNs for Image Data (28×28)

4 Loss of Spatial Structure in ANNs

5 High Number of Parameters in ANNs

6 Lack of Translation Invariance in ANNs

Artificial Neural Networks (ANNs)

Imagine predicting whether a patient has diabetes based on multiple medical features (glucose level, BMI, age, etc.). Traditional linear regression may fail to capture nonlinear interactions.

Central Question: How can we build a model that **learns nonlinear patterns** automatically?

Concept: Artificial Neural Networks

An Artificial Neural Network (ANN) is a **function approximator** inspired by the human brain. It is composed of layers of interconnected computational units called **neurons**.

The mathematical form of a single neuron is:

$$z = \sum_{i=1}^n w_i x_i + b, \quad a = f(z)$$

where: x_i = input features, w_i = weights, b = bias, $f(\cdot)$ = activation function introducing nonlinearity.

Common choices of $f(\cdot)$:

- ReLU: $f(z) = \max(0, z)$ (most hidden layers)
- Sigmoid: $f(z) = \frac{1}{1+e^{-z}}$ (binary outputs)
- Softmax: $f(z)_j = \frac{e^{z_j}}{\sum_k e^{z_k}}$ (multiclass outputs)

Artificial Neural Networks (ANNs)

Real-World Examples

- **Medical diagnosis:** Predicting diabetes, heart disease
- **Finance:** Stock price prediction, fraud detection
- **Energy:** Load forecasting in power grids
- **Computer vision:** Recognizing handwritten digits or faces

Common Mistakes

- Using too many neurons → Apply regularization (dropout, L2) to prevent overfitting
- Forgetting to normalize inputs → Scale features for stable convergence
- Using sigmoid in deep hidden layers → Prefer ReLU or its variants to avoid vanishing gradients
- Forgetting softmax for multiclass classification → Recommended to use softmax in the output layer for multiclass problems

Artificial Neural Networks (ANNs)

Caveats

- Neural networks are less interpretable than linear models
- Require large datasets and significant computational resources
- ReLU can cause “dead neurons” (weights stuck at 0)

Neurons as Light Switches with Dimmers

Think of each neuron like a light switch. Inputs are wires carrying signals. The weight decides how much current flows, the bias sets the starting brightness, and the activation function acts as a dimmer knob, deciding whether the light is dim, medium, or fully on.

Get more info → [TensorFlow Sequential API Documentation](#)

Get more info → [TensorFlow Activation Functions](#)

Summary: Artificial Neural Networks (ANNs)

Concept/Aspect	Key Explanation
Neuron formula	$z = \sum w_i x_i + b, a = f(z)$
ANN definition	Function approximator inspired by biological neurons
Role of activation	Introduces nonlinearity (ReLU, sigmoid, softmax)
Real-world examples	Diabetes prediction, fraud detection, energy load forecasting
Pitfalls	Overfitting, poor scaling, vanishing gradients, dead neurons

Current Section

- 1 Course Introduction and Marking Scheme
- 2 Artificial Neural Networks (ANNs)
- 3 Using ANNs for Image Data (28×28)
- 4 Loss of Spatial Structure in ANNs
- 5 High Number of Parameters in ANNs
- 6 Lack of Translation Invariance in ANNs

ANNs with Image Inputs

Suppose we want to classify handwritten digits (0–9) from grayscale images of size 28×28 . Each image has 784 pixels, each pixel value $\in [0, 255]$.

Central Question: How do we adapt ANNs to process image inputs effectively?

Input Representation

For an image $I \in \mathbb{R}^{28 \times 28}$:

$$x = \text{flatten}(I) \in \mathbb{R}^{784}$$

If I is the digit “5”, then:

$$x = [x_1, x_2, \dots, x_{784}]$$

where x_i are pixel intensities.

ANNs with Image Inputs

Define ANN for 28×28 Images

```
import tensorflow as tf

# Define ANN for MNIST images
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.summary()
# Expected output: Flatten -> Dense(128) -> Dense(10)
```

ANNs with Image Inputs

Numerical Example

- Image: 28×28 pixels \rightarrow 784 inputs
- ANN structure:
 - Hidden layer with 128 neurons (ReLU)
 - Output layer with 10 neurons (softmax)
- Output formula:

$$\hat{y} = \text{softmax}(W_2 \cdot f(W_1 x + b_1) + b_2)$$

where $f(\cdot)$ is ReLU activation.

Compile the Model

```
# Compile model with optimizer and loss function
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
# Expected: model is ready for training
```

ANNs with Image Inputs

Train and Evaluate ANN

```
# Load MNIST dataset
(x_train, y_train), (x_test, y_test) =
    tf.keras.datasets.mnist.load_data()

# Normalize pixel values
x_train, x_test = x_train/255.0, x_test/255.0

# Train
model.fit(x_train, y_train, epochs=5, batch_size=32, verbose=1)

# Evaluate
loss, acc = model.evaluate(x_test, y_test, verbose=0)
print("Test Accuracy:", acc) # Expected output: Accuracy ~0.97
```

Flattening Like Reading a Book

Think of the 28×28 image as a page with lines. Flattening is like reading line by line and writing all the words in a single row. The ANN only sees a long sequence, not the spatial arrangement.

[Get more info → ANN for Image Data](#)

Summary: ANNs with Image Inputs

Concept/Aspect	Key Explanation
Flattening images	Convert 28×28 matrix into 784-length vector
ANN structure	Hidden layer (ReLU) + output softmax for 10 classes
Parameter count	For 128 hidden neurons: 100,480 parameters
Best practices	Normalize pixel values, use ReLU + softmax
Limitations	Spatial info lost when flattening
TensorFlow workflow	Define → Compile → Train → Evaluate

Current Section

1 Course Introduction and Marking Scheme

2 Artificial Neural Networks (ANNs)

3 Using ANNs for Image Data (28×28)

4 Loss of Spatial Structure in ANNs

5 High Number of Parameters in ANNs

6 Lack of Translation Invariance in ANNs

Loss of Spatial Structure in ANNs

When an image is fed into a basic ANN, it is flattened into a vector of length 784 (for a 28×28 image).

Problem: Flattening destroys 2D relationships.

- A pixel from the left eye of a face may end up numerically next to a pixel from the chin.
- The model loses the fact that they were far apart in the original image.

Mathematical View

Flattening converts:

$$I \in \mathbb{R}^{28 \times 28}$$

into:

$$x = \text{flatten}(I) \in \mathbb{R}^{784}$$

The ANN processes x as:

$$h = f(Wx + b)$$

Here, W has no knowledge that x_{37} and x_{38} were neighbors in 2D, or that x_{37} and x_{742} were far apart. All distances and neighborhoods are lost.

Loss of Spatial Structure in ANNs

TensorFlow Shape Inspection

```
import tensorflow as tf

# Example input: 28x28 image
sample_input = tf.random.normal([1, 28, 28])

# After flattening
flatten_layer = tf.keras.layers.Flatten()
flattened = flatten_layer(sample_input)

print("Original shape:", sample_input.shape)
print("Flattened shape:", flattened.shape)
# Expected:
# Original shape: (1, 28, 28)
# Flattened shape: (1, 784)
```

Loss of Spatial Structure in ANNs

Numerical Example

- Pixel (10, 10) might represent part of an eye.
- Pixel (10, 11) is its right neighbor, crucial for detecting the edge of the eye.

After flattening:

$$x_{(10,10)} \mapsto x_{130}, \quad x_{(10,11)} \mapsto x_{131}$$

But pixel (11, 10) (just below the eye) becomes:

$$x_{158}$$

So the ANN treats it as unrelated, even though it is spatially connected in 2D.

Loss of Spatial Structure in ANNs

Mistakes

- Assuming ANNs can detect edges or shapes after flattening
- Believing proximity is preserved: only row-wise order is preserved, not 2D spatial locality

Warnings

- For small toy problems (like digit classification), ANNs might still work
- For complex images (faces, objects), this limitation makes ANNs ineffective

Pixels Like Scrambled Words

Imagine reading a book where all the words are taken out of order and placed in a single line. The ANN sees the “words” (pixels), but it cannot understand sentences (shapes/edges) because the order and grouping are gone.

Get more info → [CS231n: Convolutional Neural Networks for Visual Recognition](#)

Summary: Loss of Spatial Structure in ANNs

Concept/Aspect	Key Explanation
Flattening	Converts 28×28 into a 784-length vector
Loss of spatial structure	Neighbors in 2D may not stay neighbors in 1D
Effect	ANN cannot naturally detect edges, shapes, or patterns
Consequence	Poor scaling for larger/complex image tasks
TensorFlow demo	Flatten layer shows input shape (28,28) → output shape (784)

Current Section

1 Course Introduction and Marking Scheme

2 Artificial Neural Networks (ANNs)

3 Using ANNs for Image Data (28×28)

4 Loss of Spatial Structure in ANNs

5 High Number of Parameters in ANNs

6 Lack of Translation Invariance in ANNs

High Number of Parameters in ANNs

Images have thousands of pixels. In a fully connected ANN, each neuron in a layer connects to every neuron in the next.

Problem: This creates a massive number of parameters (weights + biases), which are expensive to train and prone to overfitting.

Parameter Count Formula

For a fully connected layer:

$$\# \text{params} = (\text{input_units} \times \text{output_units}) + \text{output_units}$$

where the second term corresponds to biases.

High Number of Parameters in ANNs

TensorFlow Parameter Count Demo

```
import tensorflow as tf

model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.summary()
# Expected:
# Flatten -> Dense(128): 100,480 params
# Dense(10): 1,290 params
# Total params: 101,770
```

High Number of Parameters in ANNs

Numerical Example: Two-Layer ANN

- Input: $28 \times 28 = 784$ pixels
- Hidden layer (128 neurons):

$$784 \times 128 + 128 = 100,480$$

- Output layer (10 neurons):

$$128 \times 10 + 10 = 1,290$$

- Total parameters:
- For higher resolutions (e.g., $256 \times 256 = 65,536$ inputs), parameters become astronomical.

$$100,480 + 1,290 = 101,770$$

High Number of Parameters in ANNs

Mistakes

- Ignoring parameter growth when scaling image size
- Training huge fully connected models on limited data → severe overfitting

Warnings

- Larger input images make fully connected ANNs computationally infeasible
- Even if training succeeds, such networks generalize poorly

Too Many Wires in a City

Imagine every house in a city being connected to every other house with its own wire. For a small village this might work, but for a big city the wiring becomes unmanageable. Similarly, in ANNs, as image size grows, the number of weights explodes.

Get more info → [Deep Learning Book \(Goodfellow et al.\)](#)

Summary: High Number of Parameters in ANNs

Concept/Aspect	Key Explanation
Parameter formula	$(\text{inputs} \times \text{outputs}) + \text{outputs}$
Example (28×28)	ANN with 128 hidden + 10 output neurons → 101,770 params
Problem	Parameter count grows linearly with input size
Consequence	Slow training, high memory use, overfitting risk
TensorFlow demo	<code>model.summary()</code> shows parameter explosion

Current Section

- 1 Course Introduction and Marking Scheme
- 2 Artificial Neural Networks (ANNs)
- 3 Using ANNs for Image Data (28×28)
- 4 Loss of Spatial Structure in ANNs
- 5 High Number of Parameters in ANNs
- 6 Lack of Translation Invariance in ANNs

Lack of Translation Invariance in ANNs

A standard ANN treats each input pixel as independent, tied to its fixed position.

Problem: If the object shifts in the image, the ANN cannot recognize it as the same object.

Mathematical View

Suppose $x \in \mathbb{R}^{784}$ is the flattened 28×28 image.

- Pixel (i, j) maps to x_k where $k = i \cdot 28 + j$
- If the object shifts right by 1 pixel, each input moves to a different index x_{k+1}

The ANN computes:

$$h = f(Wx + b)$$

But W is position-specific. A pattern at x_{100} is treated completely differently from the same pattern at x_{101} .

Lack of Translation Invariance in ANNs

TensorFlow: Shifted Inputs Demonstration

```
import tensorflow as tf

# Simulate a "digit" as a 1D tensor with a peak
x_center = tf.concat([tf.zeros(5), tf.ones(3), tf.zeros(20)], axis=0)
x_shifted = tf.concat([tf.zeros(6), tf.ones(3), tf.zeros(19)], axis=0)

print("Center index of peak:", tf.argmax(x_center).numpy())
print("Shifted index of peak:", tf.argmax(x_shifted).numpy())
# Expected:
# Center index of peak: 5
# Shifted index of peak: 6
```

Lack of Translation Invariance in ANNs

Numerical Example

- Digit “7” centered → activates pixels x_{300} to x_{330}
- Shifted right → activates x_{301} to x_{331}
- ANN must learn two different weight sets for the same digit

If there are k possible shifts, the ANN must effectively memorize k different versions of the same object.

Mistakes

- Assuming a trained ANN can recognize objects anywhere in the image
- Believing deeper layers automatically solve translation invariance

Lack of Translation Invariance in ANNs

Warnings

- Lack of translation invariance leads to data inefficiency
- The network must see the same object in many positions during training
- Especially harmful for larger images where objects can appear anywhere

Recognizing Faces Only at One Spot

Imagine teaching someone to recognize a face but only if it is printed in the center of a page. If you shift the photo slightly left, they claim it is a new, unfamiliar image. That is how ANNs treat translations.

Get more info → [Convolutions for Images and Translation Invariance](#)

Summary: Lack of Translation Invariance in ANNs

Concept/Aspect	Key Explanation
Translation invariance	ANN does not naturally recognize shifted objects
Mathematical reason	Different pixel indices → different weights
Effect	Network must learn multiple redundant patterns
Consequence	Poor generalization and inefficiency
TensorFlow demo	Shifted inputs show index changes of activations